A Performance Analysis of Differing Data Structures

Vincent Curran

CSCI 2270-104

April 27th, 2020

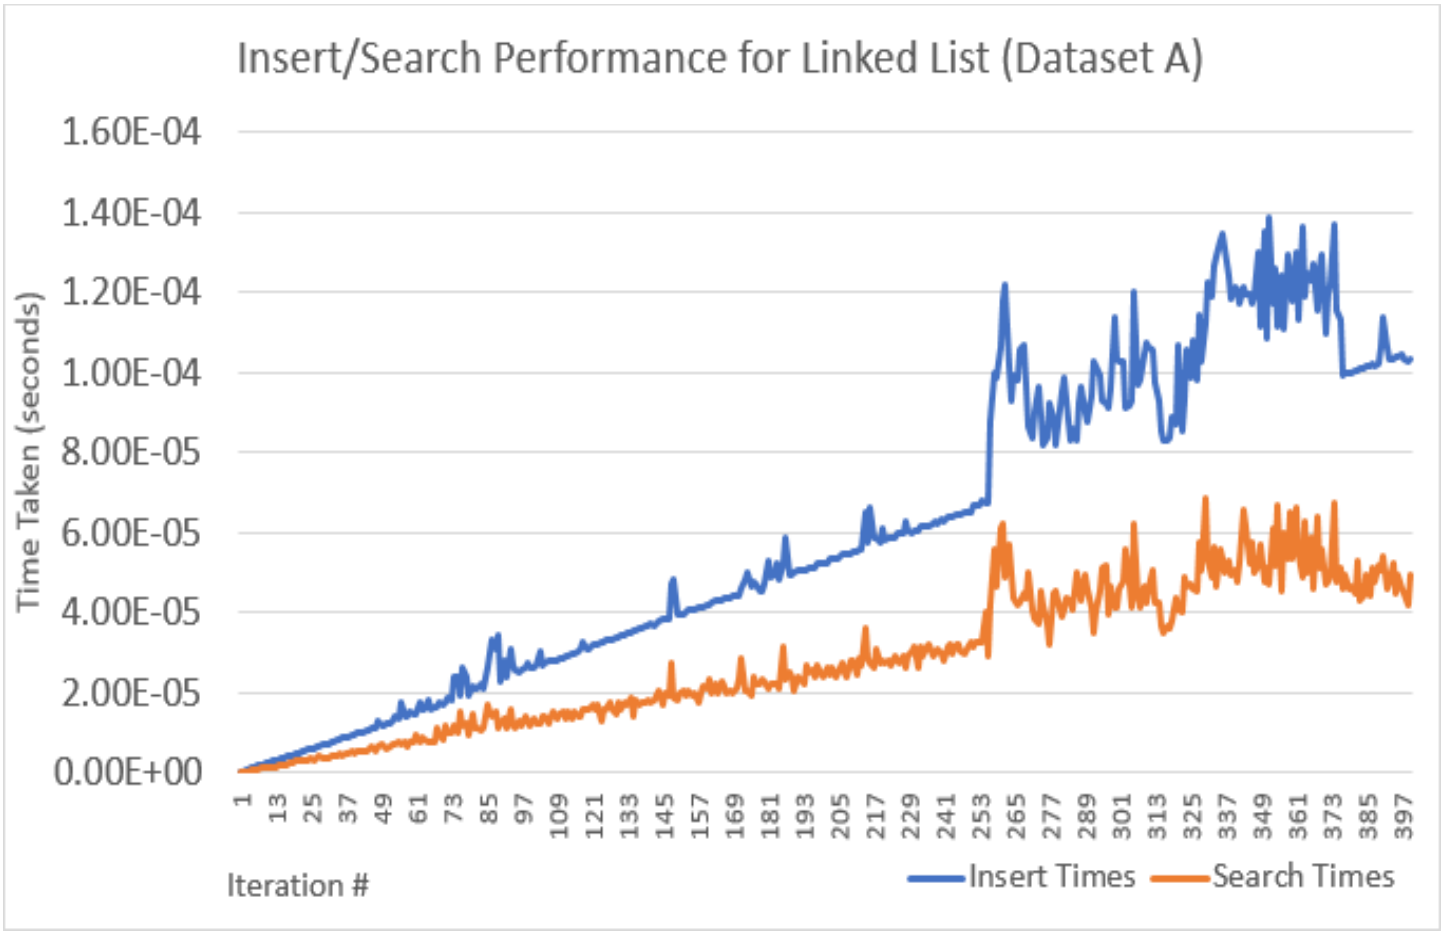
Insert/Search Performance for Linked List (Dataset A)
Time Taken (seconds)
1.60E-04
1.40E-04
1.20E-04
1.00E-04
8.00E-05
6.00E-05
4.00E-05
2.00E-05
0.00E+00
1 13 25 37 49 61 73 85 97 109 121 133 145 157 169 181 193 205 217 229 241 253 265 277 289 301 313 325 337 349 361 373 385 397
Iteration #
Insert Times
Search Times

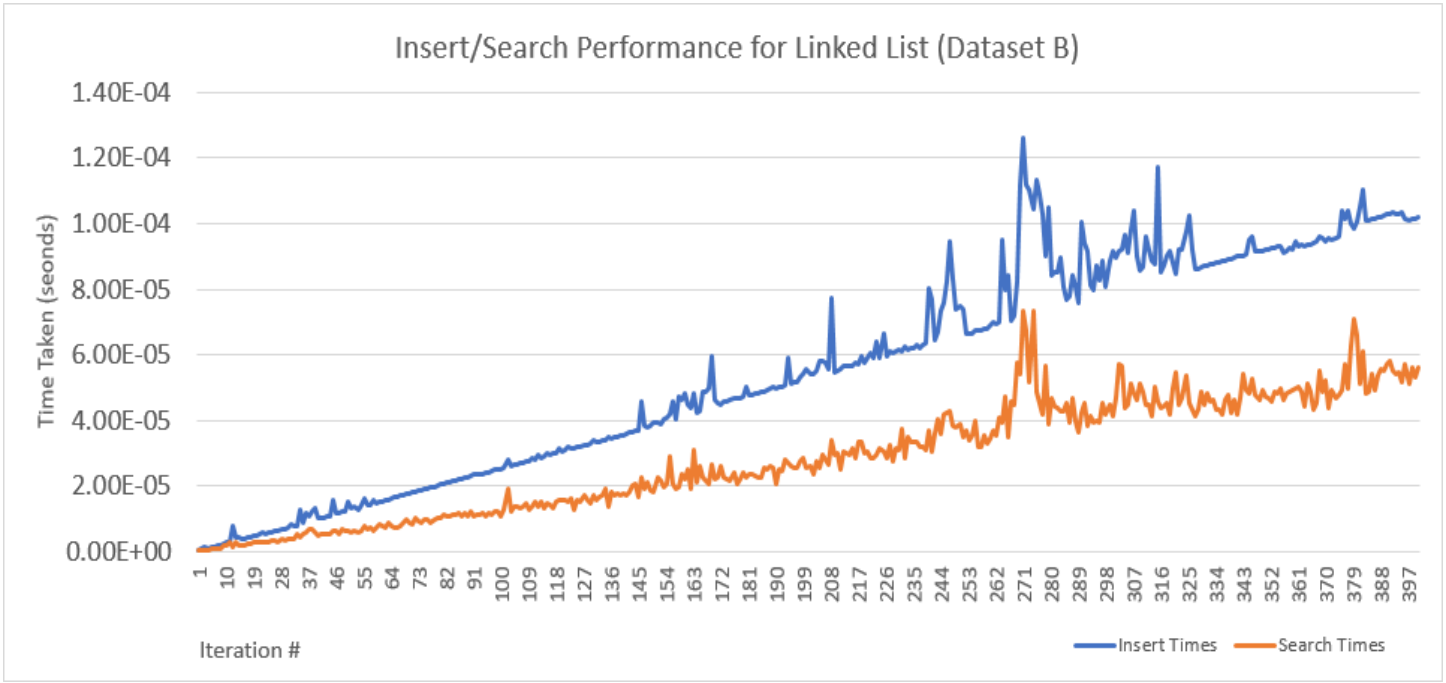
Insert/Search Performance for Linked List (Dataset B)
Time Taken (seonds)
1.40E-04
1.20E-04
1.00E-04
8.00E-05
6.00E-05
4.00E-05
2.00E-05
0.00E+00
1 10 19 28 37 46 55 64 73 82 91 100 109 118 127 136 145 154 163 172 181 190 199 208 217 226 235 244 253 262 271 280 289 298 307 316 325 334 343 352 361 370 379 388 397
Iteration #
Insert Times
Search Times

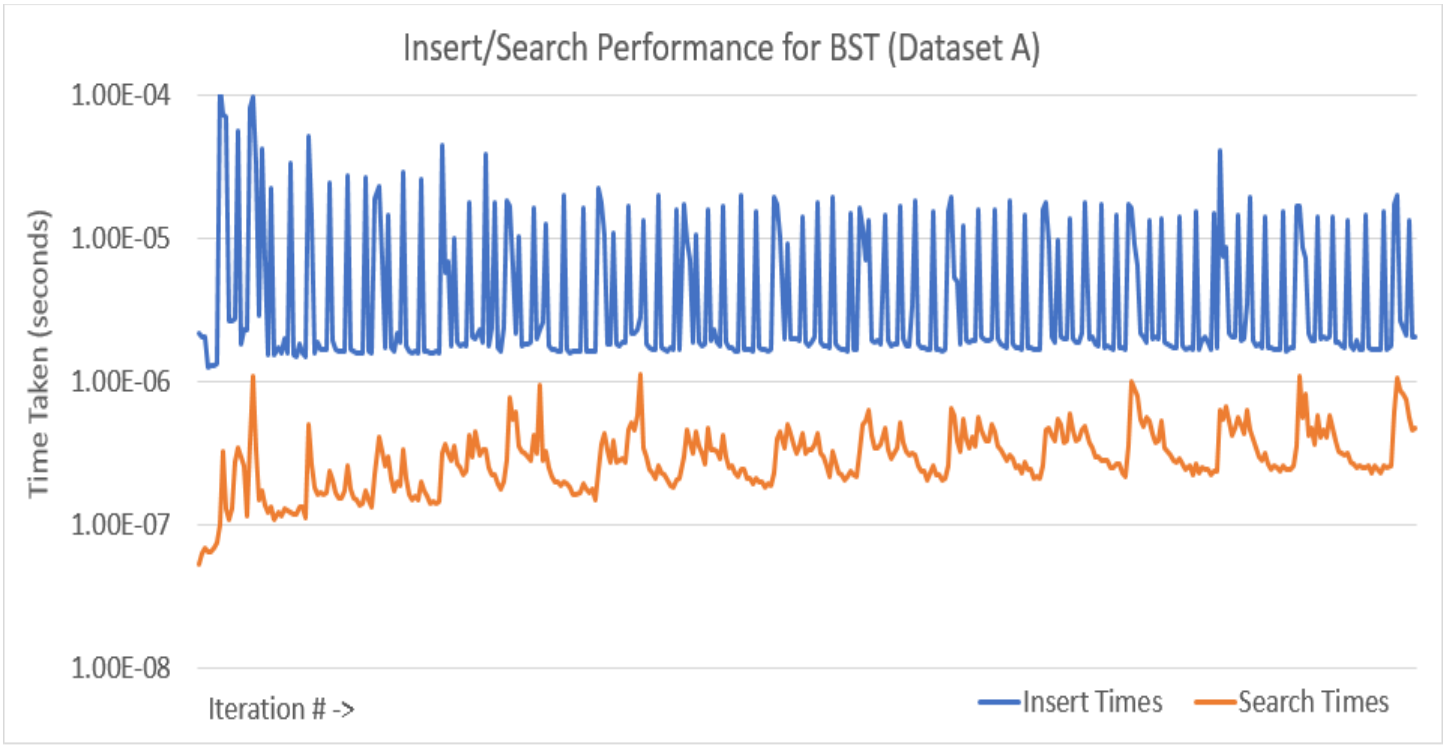
Insert/Search Performance for BST (Dataset A)
1.00E-04
1.00E-05
1.00E-06
1.00E-07
1.00E-08
Time Taken (seconds)
Iteration # ->
Insert Times
Search Times

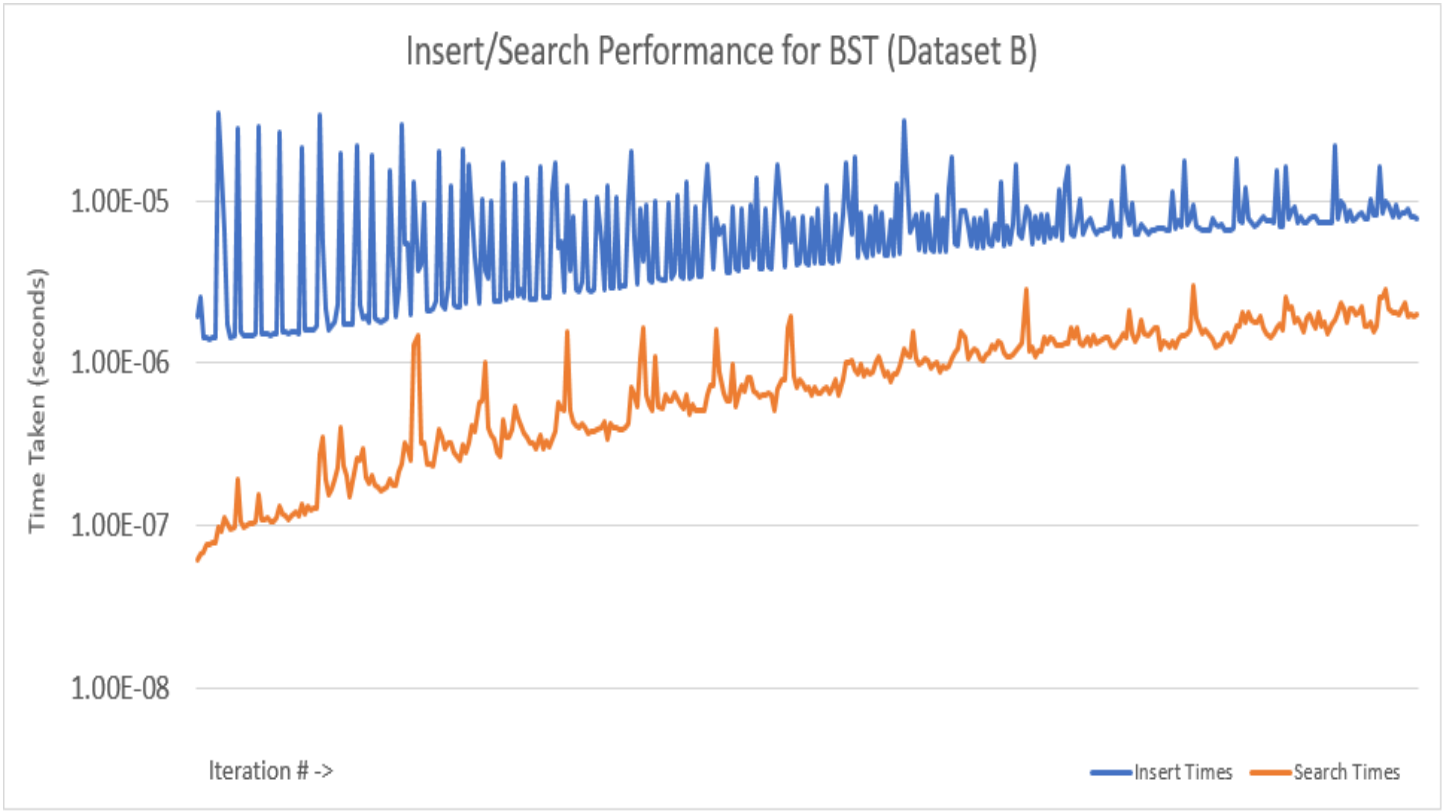
Insert/Search Performance for BST (Dataset B)
1.00E-05
1.00E-06
1.00E-07
1.00E-08
Time Taken (seconds)
Iteration # ->
Insert Times
Search Times

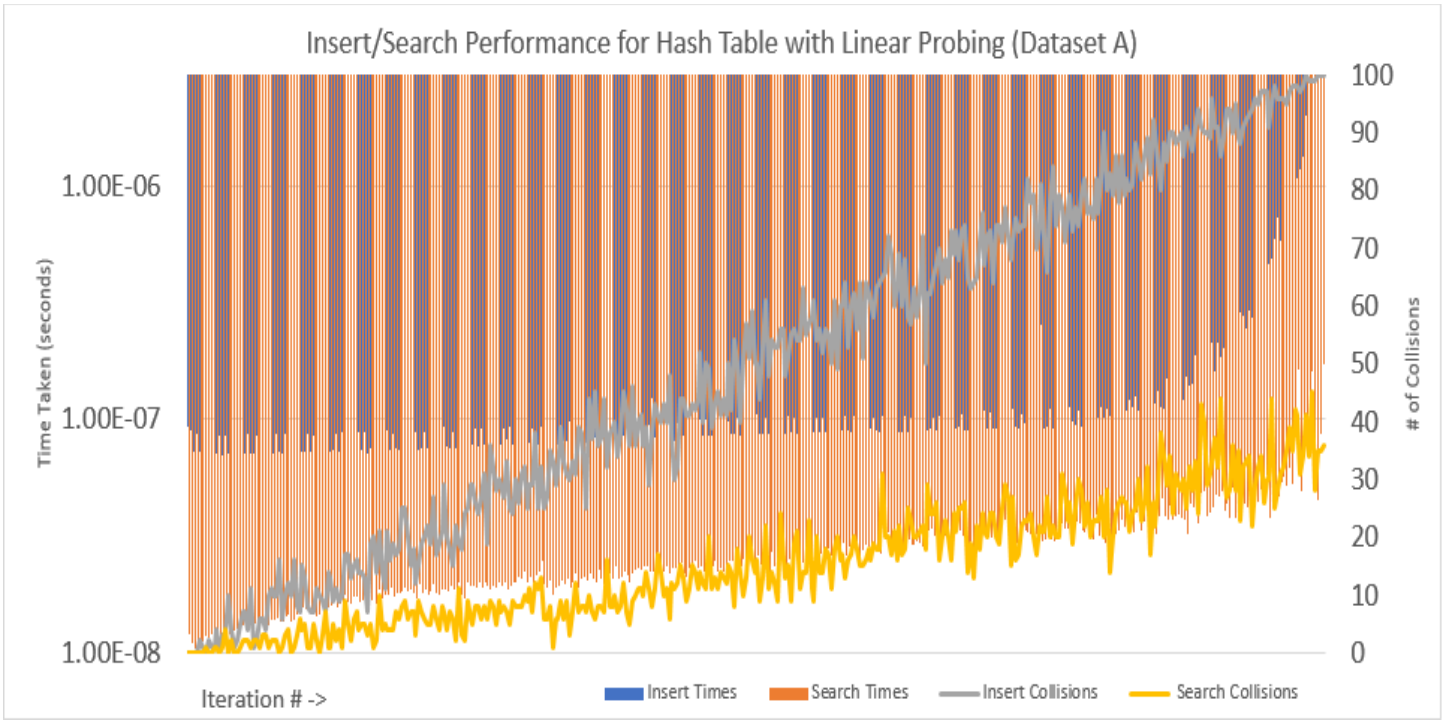
Insert/Search Performance for Hash Table with Linear Probing (Dataset A)
Time Taken (seconds)
1.00E-06
1.00E-07
1.00E-08
# of Collisions
100
90
80
70
60
50
40
30
20
10
0
Iteration # ->
Insert Times
Search Times
Insert Collisions
Search Collisions

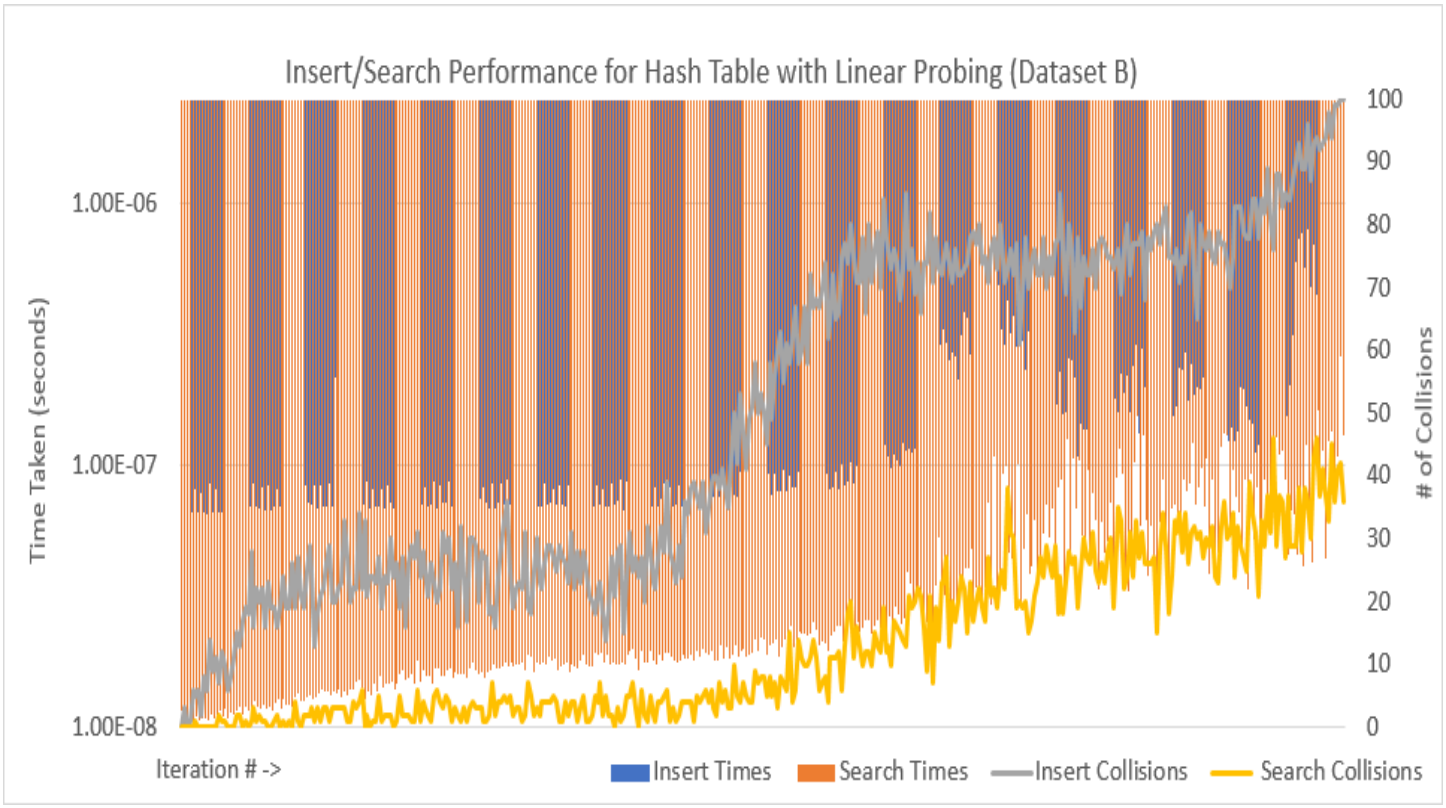
Insert/Search Performance for Hash Table with Linear Probing (Dataset B)
Time Taken (seconds)
1.00E-06
1.00E-07
1.00E-08
# of Collisions
100
90
80
70
60
50
40
30
20
10
0
Iteration # ->
Insert Times
Search Times
Insert Collisions
Search Collisions

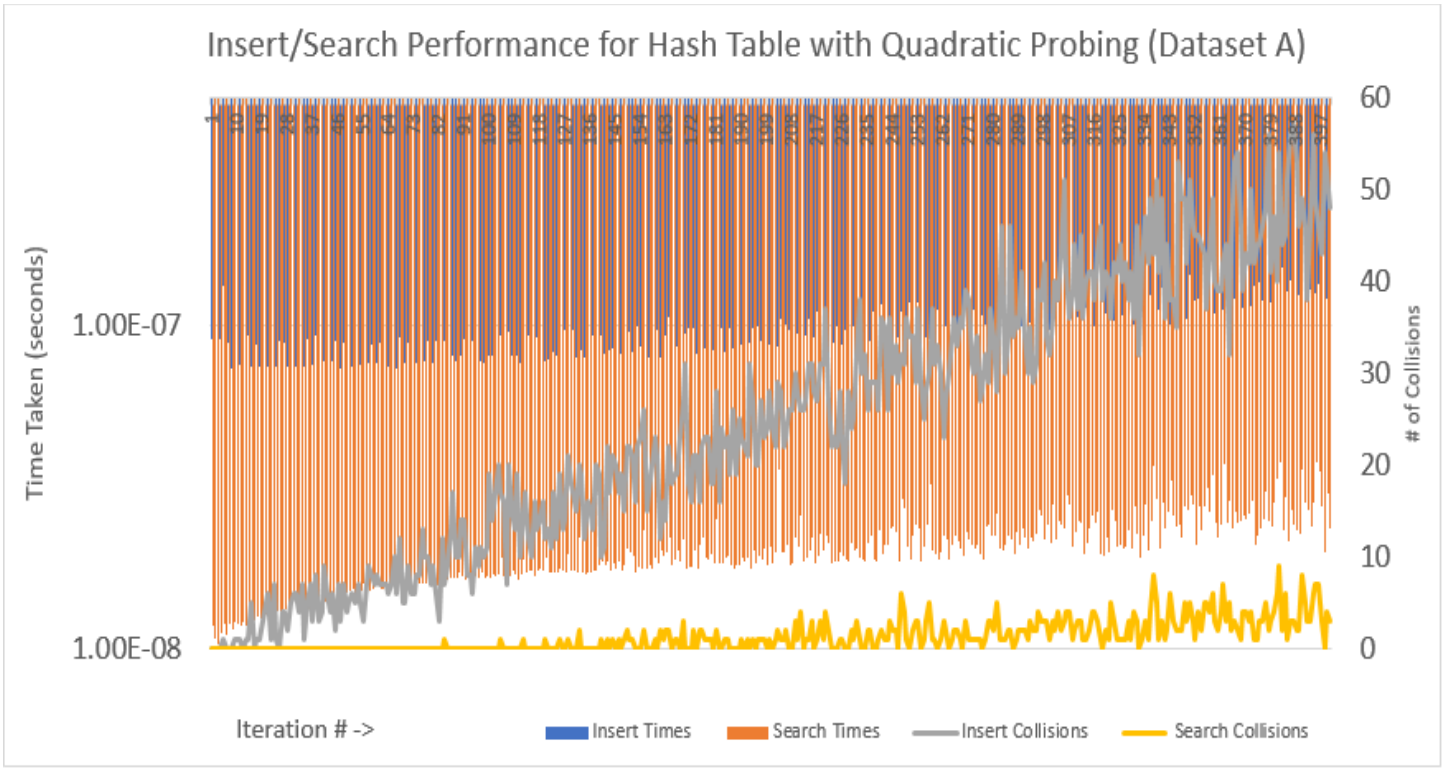
Insert/Search Performance for Hash Table with Quadratic Probing (Dataset A)
Time Taken (seconds)
1.00E-07
1.00E-08
# of Collisions
60
50
40
30
20
10
0
Iteration # ->
Insert Times
Search Times
Insert Collisions
Search Collisions

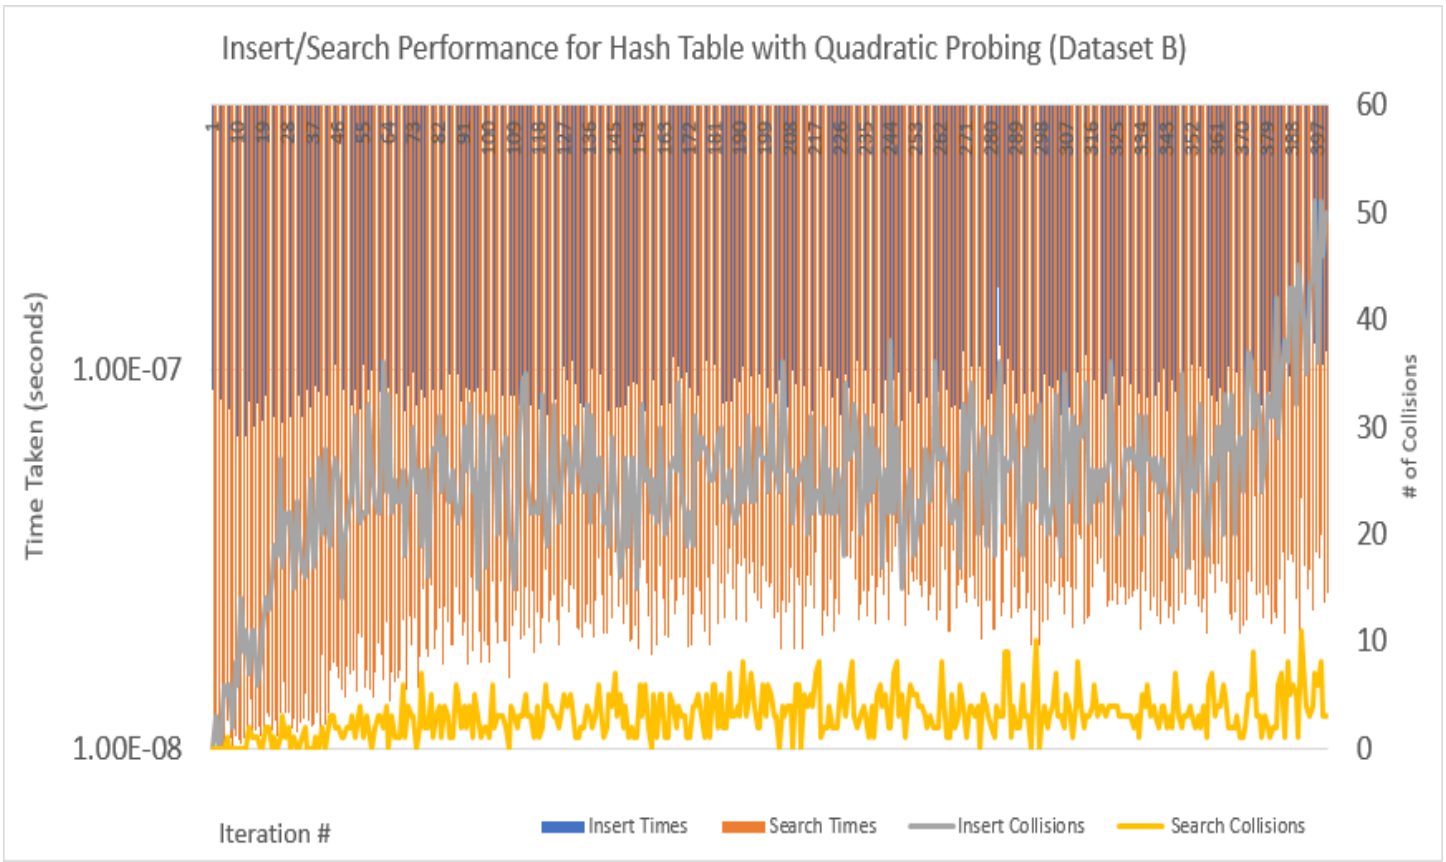
Insert/Search Performance for Hash Table with Quadratic Probing (Dataset B)
Time Taken (seconds)
1.00E-07
1.00E-08
# of Collisions
60
50
40
30
20
10
0
Iteration #
Insert Times
Search Times
Insert Collisions
Search Collisions

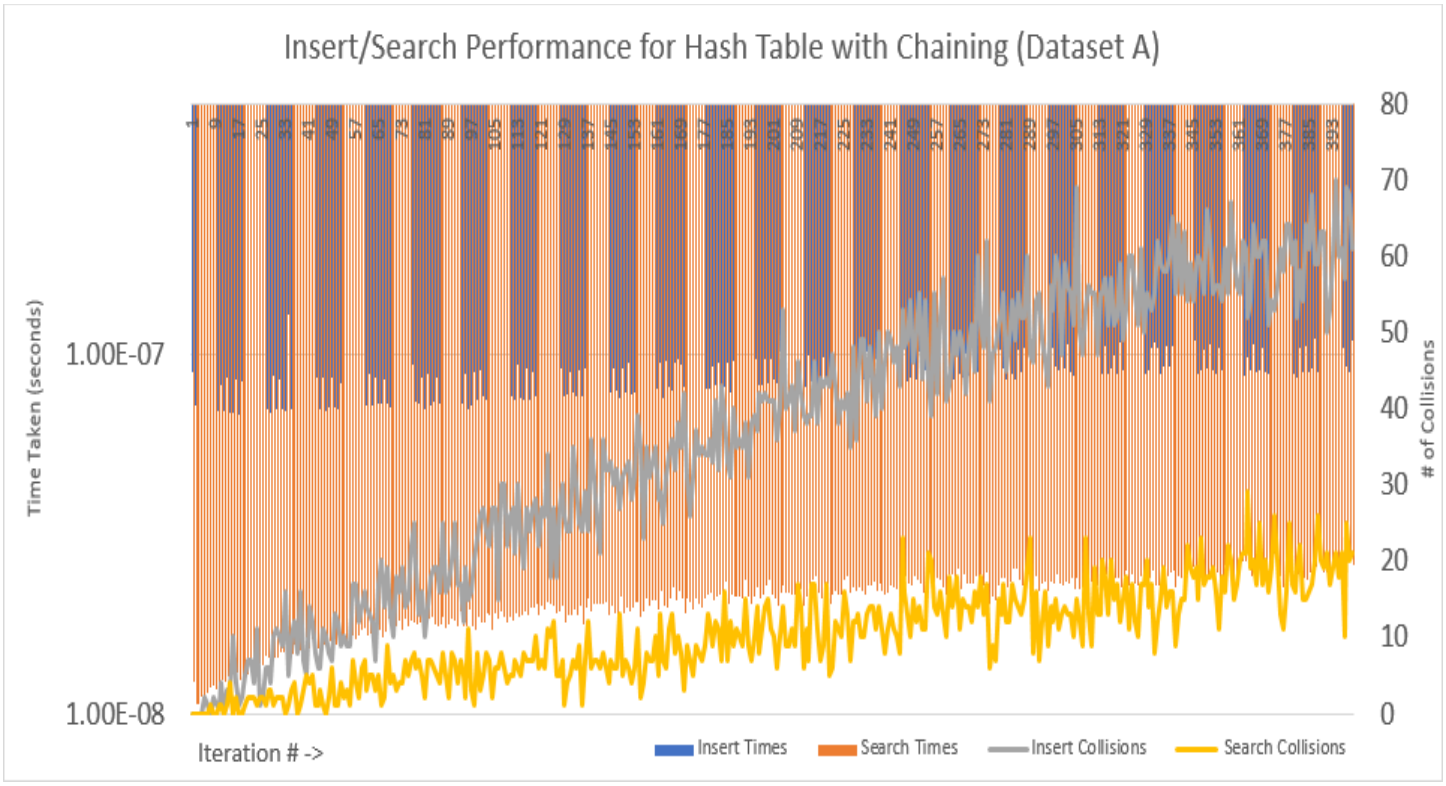
Insert/Search Performance for Hash Table with Chaining (Dataset A)
Time Taken (seconds)
1.00E-07
1.00E-08
# of Collisions
80
70
60
50
40
30
20
10
0
Iteration # ->
Insert Times
Search Times
Insert Collisions
Search Collisions

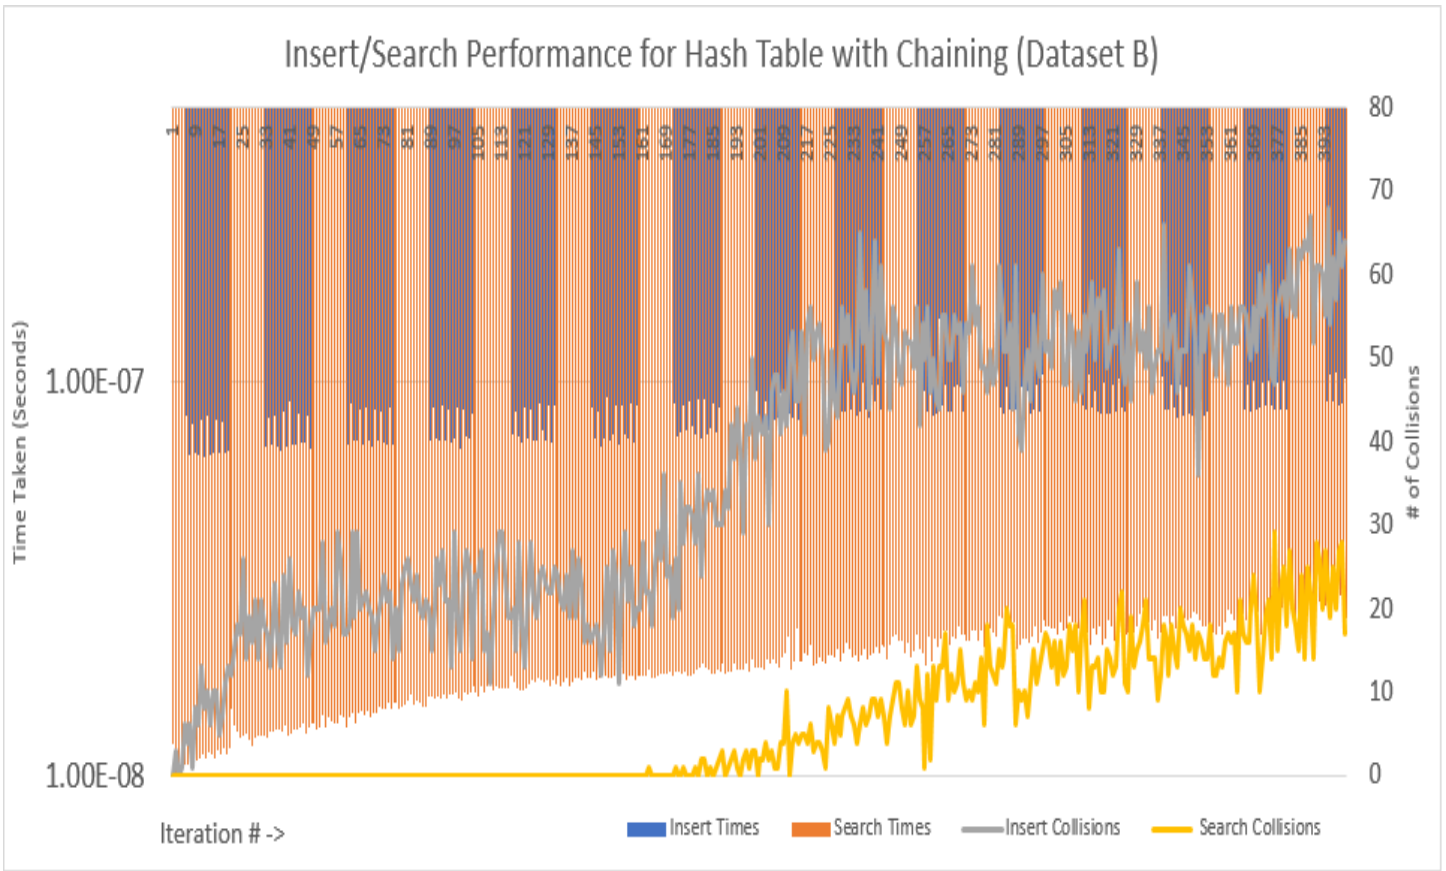
Insert/Search Performance for Hash Table with Chaining (Dataset B)
Time Taken (Seconds)
1.00E-07
1.00E-08
# of Collisions
80
70
60
50
40
30
20
10
0
Iteration # ->
Insert Times
Search Times
Insert Collisions
Search Collisions

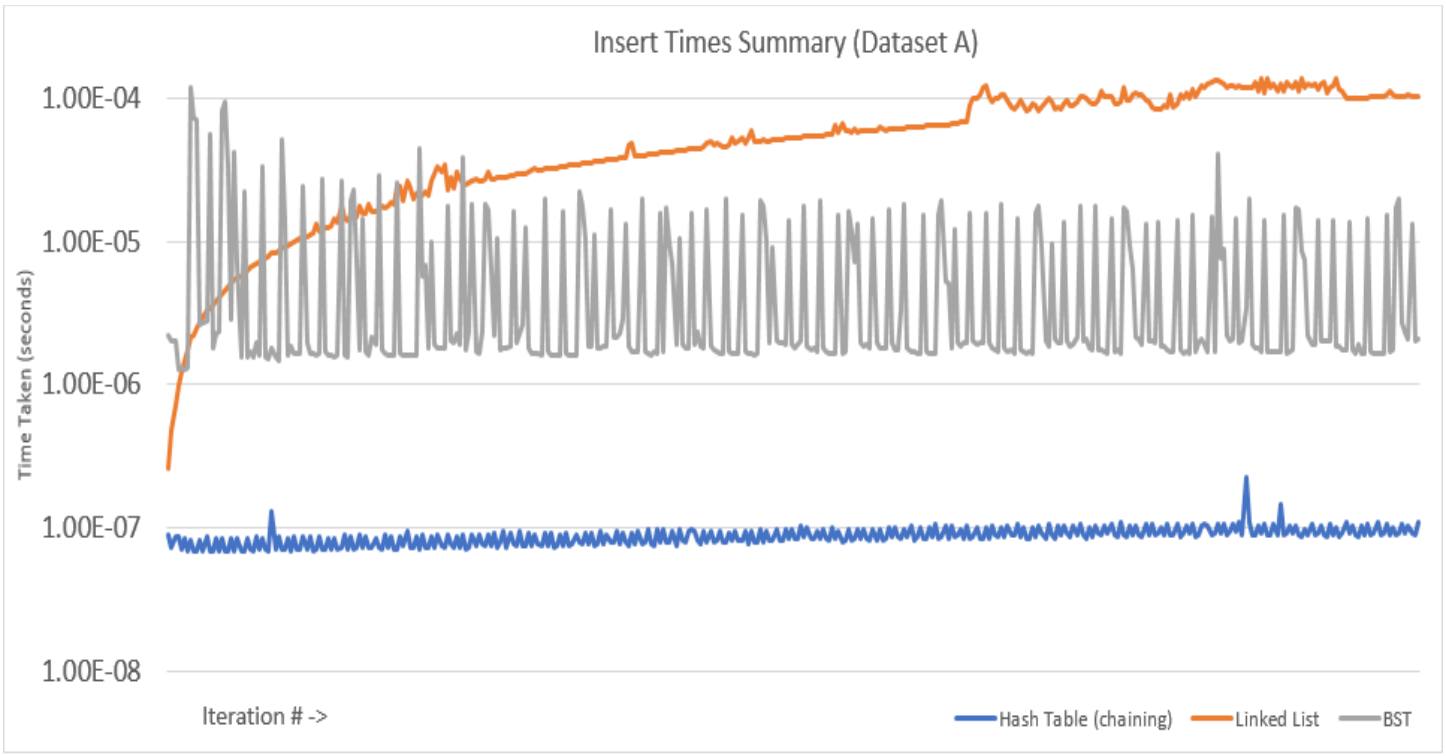
Insert Times Summary (Dataset A)
1.00E-04
1.00E-05
1.00E-06
1.00E-07
1.00E-08
Time Taken (seconds)
Iteration # ->
Hash Table (chaining)
Linked List
BST

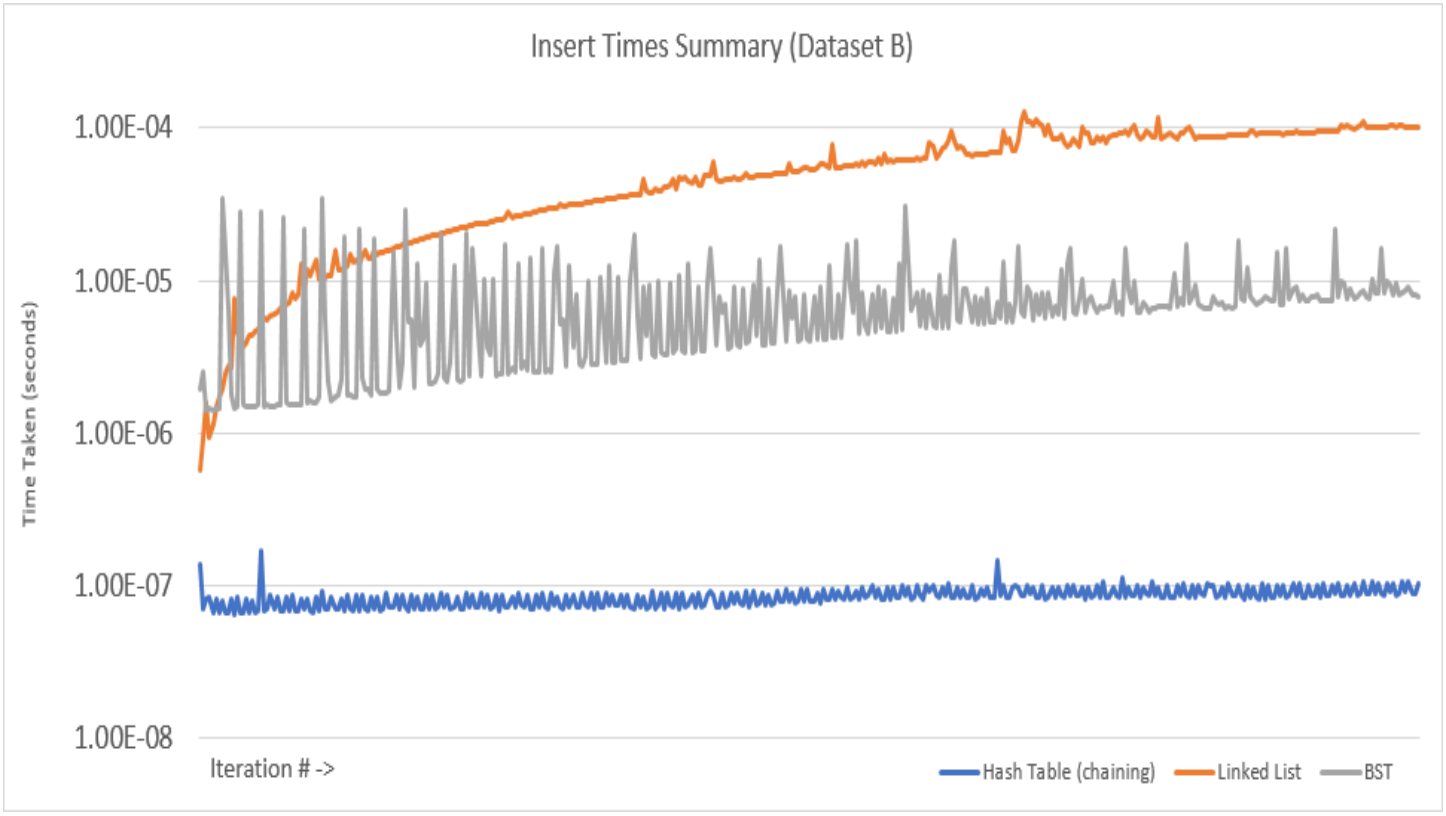
Insert Times Summary (Dataset B)
1.00E-04
1.00E-05
1.00E-06
1.00E-07
1.00E-08
Time Taken (seconds)
Iteration # ->
Hash Table (chaining)
Linked List
BST

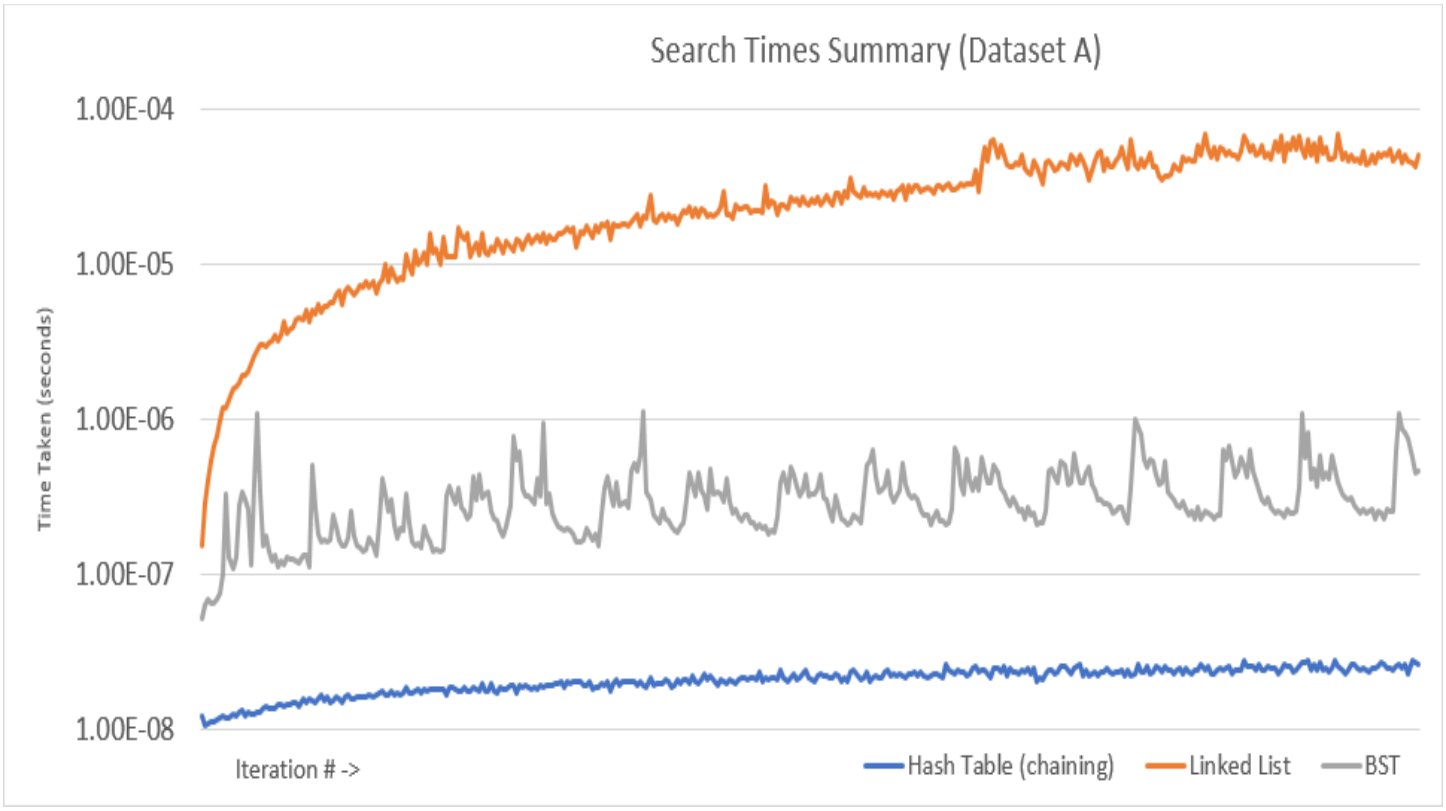
Search Times Summary (Dataset A)
1.00E-04
1.00E-05
1.00E-06
1.00E-07
1.00E-08
Time Taken (seconds)
Iteration # ->
Hash Table (chaining)
Linked List
BST

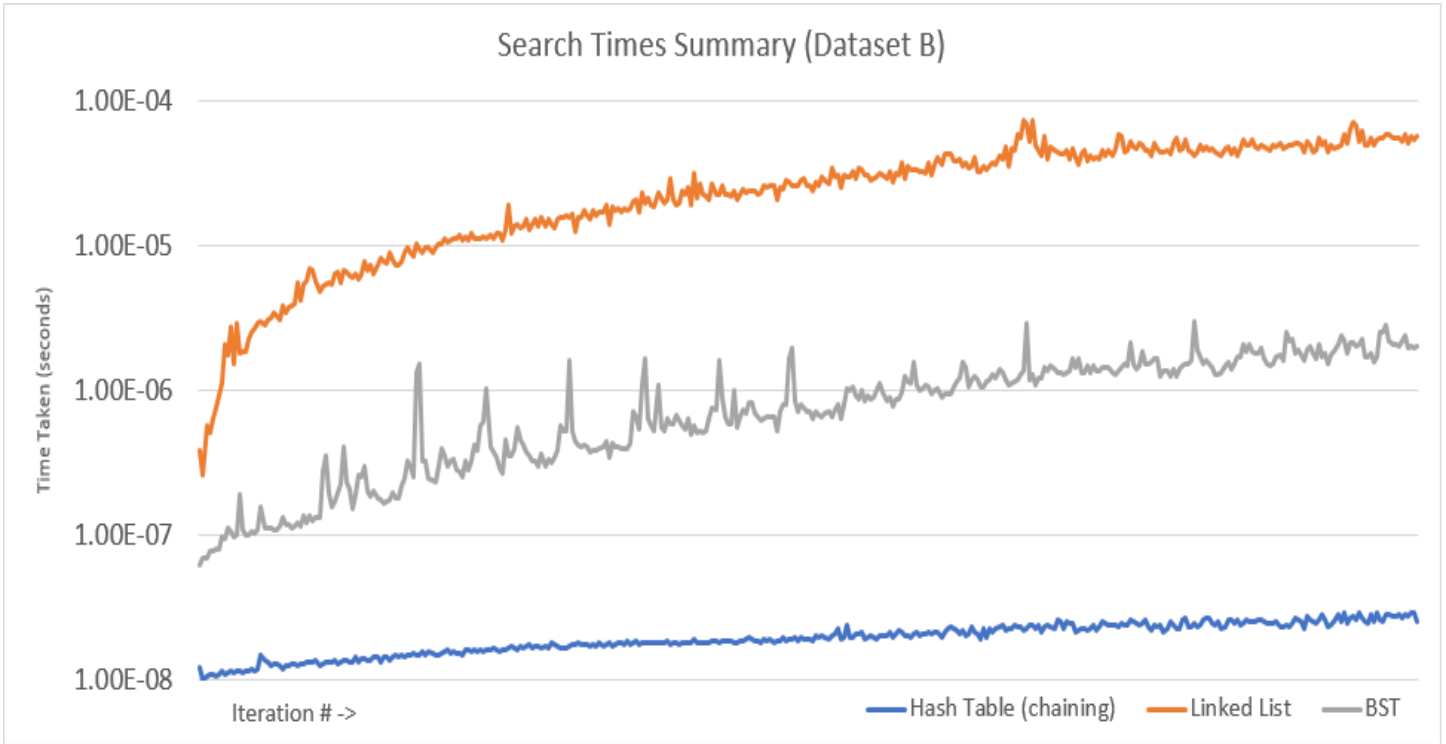
Search Times Summary (Dataset B)
1.00E-04
1.00E-05
1.00E-06
1.00E-07
1.00E-08
Time Taken (seconds)
Iteration # ->
Hash Table (chaining)
Linked List
BST

After conducting the necessary tests and analyzing the results, I have concluded the best data structure for the USPS to use for tracking shipping numbers is a hash table utilizing chaining for collision resolution. The linked list was the first and easiest to rule out, as it performed the slowest to start and got much worse as the data scaled up. This is because to find or insert anything, you must traverse past every other point already in the list. The binary search tree was next to go, and while it performed better than the linked list, it was subject to the same issues with data scaling as the list. This makes sense, because the binary search tree is effectively a method of cutting through the list, so there's less datapoints to go through before reaching the desired place. I then found that a hash table which uses chaining, linear probing, or quadratic probing for collision resolution can both insert and search through their data several orders of magnitude faster than when compared to a linked list or even binary search tree. This is because the hash table reduces the time it takes to reach any data point to a single line, then if a collision occurs it's likely only a few more operations. Among the hash tables, the quadratic probing and chaining methods of collision resolution proved superior to the linear probing method. This was due to a higher number of collisions during linear probing, which led to a large increase in both insert and search functions when the table neared capacity. I ultimately chose to recommend the hash table with chaining as the preferred method of data storage for the USPS because to get the quadratic probing method to have the same performance, it needed a table size twice as large (the quadratic method cannot guarantee a spot in the table if table size is equal to what was used for linear probing). Also, while it proved to have less collisions than with the chaining method, the latter seemed to solve the collisions faster as it did not seem to have much of an impact on time performance.